

softmax

February 11, 2020

1 Softmax Classifier

This exercise guides you through the process of classifying images using a Softmax classifier. As part of this you will:

- Implement a fully vectorized loss function for the Softmax classifier
- Calculate the analytical gradient using vectorized code
- Tune hyperparameters on a validation set
- Optimize the loss function with Stochastic Gradient Descent (SGD)
- Visualize the learned weights

```
[1]: # start-up code!
import random

import matplotlib.pyplot as plt
import numpy as np

%matplotlib inline
plt.rcParams['figure.figsize'] = (10.0, 8.0) # set default size of plots
plt.rcParams['image.interpolation'] = 'nearest'
plt.rcParams['image.cmap'] = 'gray'

# for auto-reloading extenrnal modules
# see http://stackoverflow.com/questions/1907993/
  ↳ autoreload-of-modules-in-ipython
%load_ext autoreload
%autoreload 2
```

```
[2]: from load_cifar10_tvt import load_cifar10_train_val

X_train, y_train, X_val, y_val, X_test, y_test = load_cifar10_train_val()
print("Train data shape: ", X_train.shape)
print("Train labels shape: ", y_train.shape)
print("Val data shape: ", X_val.shape)
print("Val labels shape: ", y_val.shape)
print("Test data shape: ", X_test.shape)
print("Test labels shape: ", y_test.shape)
```

Train, validation and testing sets have been created as

X_i and y_i where $i=\text{train, val, test}$

Train data shape: (3073, 49000)

Train labels shape: (49000,)

Val data shape: (3073, 1000)

Val labels shape: (1000,)

Test data shape: (3073, 1000)

Test labels shape: (1000,)

Code for this section is to be written in `cs231n/classifiers/softmax.py`

```
[17]: # Now, implement the vectorized version in softmax_loss_vectorized.

import time

from cs231n.classifiers.softmax import softmax_loss_vectorized

# gradient check.
from cs231n.gradient_check import grad_check_sparse

W = np.random.randn(10, 3073) * 0.0001

tic = time.time()
loss, grad = softmax_loss_vectorized(W, X_train, y_train, 0.00001)
toc = time.time()
print("vectorized loss: %e computed in %fs" % (loss, toc - tic))

# As a rough sanity check, our loss should be something close to -log(0.1).
print("loss: %f" % loss)
print("sanity check: %f" % (-np.log(0.1)))

f = lambda w: softmax_loss_vectorized(w, X_train, y_train, 0.0)[0]
grad_numerical = grad_check_sparse(f, W, grad, 10)
```

vectorized loss: 2.371848e+00 computed in 0.549096s

loss: 2.371848

sanity check: 2.302585

numerical: 3.938143 analytic: 3.938142, relative error: 8.559665e-09

numerical: -0.214155 analytic: -0.214155, relative error: 3.923481e-08

numerical: -0.999146 analytic: -0.999146, relative error: 8.871554e-09

numerical: -0.130245 analytic: -0.130245, relative error: 1.684215e-07

numerical: 1.668923 analytic: 1.668923, relative error: 2.392686e-09

numerical: -0.798887 analytic: -0.798887, relative error: 4.165373e-08

numerical: 0.565113 analytic: 0.565113, relative error: 2.727624e-08

numerical: -0.977033 analytic: -0.977033, relative error: 2.373019e-08

numerical: -1.104690 analytic: -1.104690, relative error: 4.872252e-08

numerical: 2.468833 analytic: 2.468833, relative error: 9.964113e-09


Code for this section is to be written in `cs231n/classifiers/linear_classifier.py`

```
[36]: # Now that efficient implementations to calculate loss function and gradient of
      ↪ the softmax are ready,
      # use it to train the classifier on the cifar-10 data
      # Complete the `train` function in cs231n/classifiers/linear_classifier.py

      from cs231n.classifiers.linear_classifier import Softmax

      classifier = Softmax()
      loss_hist = classifier.train(
          X_train,
          y_train,
          learning_rate=1e-4,
          reg=1e-5,
          num_iters=100,
          batch_size=200,
          verbose=False,
      )
      # Plot loss vs. iterations
      plt.plot(loss_hist)
      plt.xlabel("Iteration number")
      plt.ylabel("Loss value")
```

```
[36]: Text(0, 0.5, 'Loss value')
```



```
[37]: # Complete the `predict` function in cs231n/classifiers/linear_classifier.py
# Evaluate on test set
y_test_pred = classifier.predict(X_test)
test_accuracy = np.mean(y_test == y_test_pred)
print("softmax on raw pixels final test set accuracy: %f" % (test_accuracy,))
```

softmax on raw pixels final test set accuracy: 0.253000

```
[35]: # Visualize the learned weights for each class
w = np.array(classifier.W[:, :-1]) # strip out the bias
w = w.reshape(10, 32, 32, 3)

w_min, w_max = np.min(w), np.max(w)

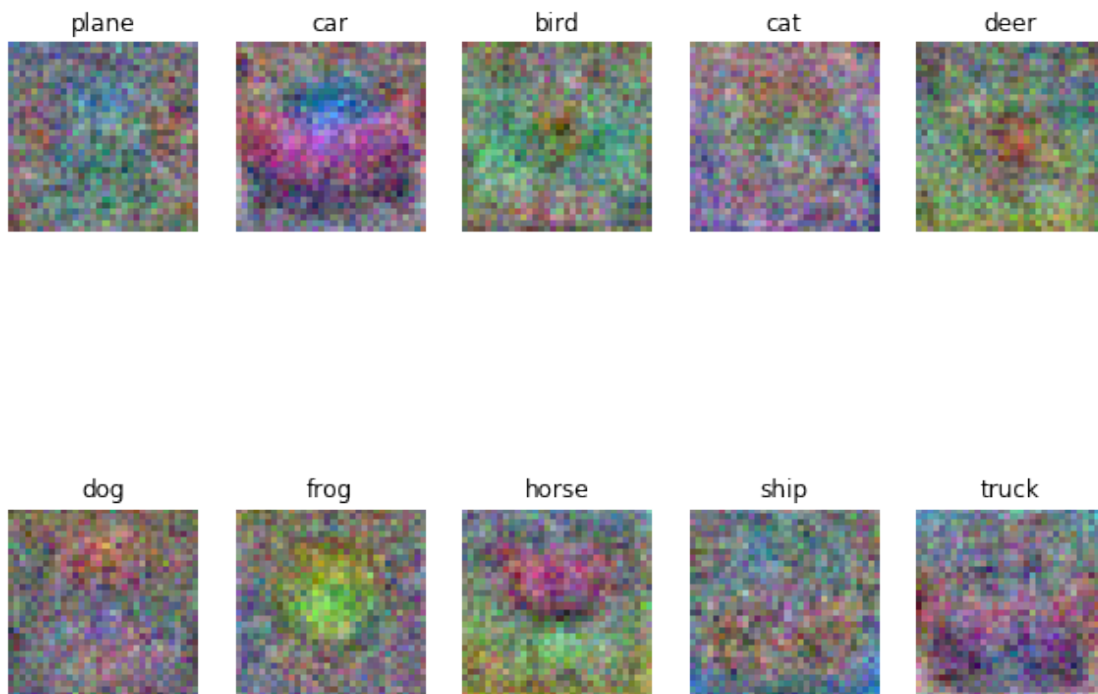
classes = [
    "plane",
    "car",
    "bird",
    "cat",
```

```

"deer",
"dog",
"frog",
"horse",
"ship",
"truck",
]
for i in range(10):
    plt.subplot(2, 5, i + 1)

    # Rescale the weights to be between 0 and 255
    wimg = 255.0 * (w[i].squeeze() - w_min) / (w_max - w_min)
    plt.imshow(wimg.astype("uint8"))
    plt.axis("off")
    plt.title(classes[i])

```



[]:

[]: